

DATA CARD

Input vs. Output	The model receives text input: a natural-language prompt (e.g., "Tell a moral market tale about returning a lost cowrie shell") combined with structured metadata: a theme and <i>culture_region</i> tag that condition generation, plus an internally selected "Style cue" (a short excerpt from a matched reference document, chosen via TF-IDF similarity) that provides stylistic grounding. The output is a short, original English-language folktale (a Story string), typically 30–70 words, targeting the training corpus's compressed narrative structure (setup → turn → consequence → proverb-like closing → lasting change).
Classification problem	This is primarily a text-generation task, not a classification task. However, generation is conditioned on two categorical fields provided by the competition dataset schema: theme (6 classes: <i>trickster</i> , <i>origin_myth</i> , <i>moral_tale</i> , <i>hero_journey</i> , <i>animal_fable</i> , <i>community_wisdom</i>) and <i>culture_region</i> (5 classes: <i>west_africa</i> , <i>east_africa</i> , <i>southern_africa</i> , <i>central_africa</i> , <i>diaspora</i>). These categories were defined by the competition organizers as part of the dataset, not chosen by our team; we use them both to select relevant stylistic context per prompt and to check regional/thematic balance in our evaluation.
Labeling Strategy	Training targets (the reference stories the model learns to produce) come from the competition's own provided synthetic corpus, which the organizers describe explicitly as "educational synthetic text," not transcribed oral tradition or human-labeled data. For our supplemental documents, theme and region tags were assigned by our team based on the real-world tradition each document was grounded in (e.g., a Zulu-tradition motif tagged <i>southern_africa</i>). We checked labeling consistency by cross-referencing each supplemental document's claimed source against live, independent verification: this process actually caught and corrected two inaccurate citations during our own review, which we treat as evidence the check works rather than something to hide.
Data creation	Our primary corpus is provided directly by the competition organizers (<i>documents.csv</i> , <i>train_prompts.csv</i> , <i>test_prompts.csv</i>), licensed CC-BY-4.0 (package)/CC0-1.0 (synthetic documents). We supplemented this with 16 original documents we authored ourselves: 6 are prose retellings grounded in specific, independently verified traditional African motifs (not copied text; original composition inspired by well-documented, community-attributed oral traditions), and 10 are original compositions in a similar stylistic register without a claimed specific source. We additionally hand-compressed 6 of the verified-motif documents into new training examples in the corpus's house style. No data was scraped in bulk from the web, and no crowd-sourced or third-party labeled datasets were used.
Web collection ethics	We did not scrape or bulk-collect training text from the web. Web search was used only to identify and verify that a given traditional motif (e.g., "Anansi and the Sky God's stories," "the Moon and Hare's message") is genuinely well-documented and community-attributed, after which we wrote fully original prose ourselves rather than reproducing any source's specific wording. This avoids both copyright issues and the risk of misattributing invented content to a real source (an error we caught and corrected during our own review). We avoided any motif that reads as culturally restricted or ceremonial, restricting ourselves to material already openly and widely published/retold. No personal, private, or consent-requiring data was collected at any point.
Bias and representation	The organizer-provided corpus has its weakest coverage for the "diaspora" region (fewest source documents); notably, the competition's actual test set contains no diaspora-themed prompts, so this gap doesn't affect our leaderboard evaluation but remains a real limitation for broader real-world use. Given time and verification constraints, our own supplemental documents also could not achieve fully even coverage across all 5 regions × 6 themes. Some combinations still rely solely on the original synthetic corpus. Because the underlying training data is synthetic rather than sourced from real community members, no group's authentic lived voice is directly represented in training; we disclose this explicitly rather than presenting the dataset as more representative than it is.

Values and community alignment	This dataset design reflects the values named in our Challenge 1 problem statement — cultural authenticity, inclusivity, transparency, and methodological rigor. Concretely: we disclosed all supplemental data sources on the competition forum (transparency); attempted balanced coverage across regions and themes rather than defaulting to whichever sources were easiest to find (inclusivity); used original retellings grounded in verified motifs rather than exploitative scraping (respect/non-appropriation); and explicitly documented the synthetic nature of our primary training data rather than overclaiming authenticity (safety against misrepresentation). Community benefit and accountability are addressed by framing the resulting model as an early-stage educational prototype, not an authoritative source — every output carries a disclaimer, and our sourcing methodology is fully disclosed for public scrutiny rather than presented as beyond question.
---------------------------------------	---