

IMPACT STATEMENT

Purpose of the project: Our project fine-tunes a small language model (Gemma-2-2b-it, via LoRA) to generate original English-language folktales that are structurally and stylistically authentic to African oral storytelling traditions, spanning the continent's regions and diaspora. It addresses the risk that AI systems, trained mostly on Western literary text, default to Western narrative conventions even when asked to produce "African-themed" content, flattening the distinctive proverb-driven, trickster-centered structures of African oral tradition.

Positive impacts: If developed further, a model like this could support educators and cultural organizations looking for supplementary creative material that reflects African narrative structure rather than generic fairy-tale conventions. It also demonstrates, concretely, how a small team can approach cultural representation in AI responsibly, through verified sourcing and honest disclosure, offering a model for other student or community projects working with limited data and compute.

Potential risks and harms: The clearest risk is overclaiming authenticity. Our primary training corpus is synthetic, not transcribed oral tradition, and even our verified supplemental sources are original retellings, not direct community transmission. If presented without that context, outputs could be mistaken for authoritative cultural material. A second risk is uneven regional representation, since some regions, particularly the diaspora, have thinner source coverage than others. There is also a general risk shared by all generative AI: outputs could subtly reinforce stereotype if training data leaned on exoticized framing, though we screened against this deliberately.

Values: The project is built around cultural authenticity, inclusivity, transparency, and community well-being. We prioritized honest representation over polish, disclosed every external source, and designed outputs to carry disclaimers rather than imply false authority.

Mitigation strategies: We verified supplemental sources against independent references before use, correcting inaccuracies we found during our own review. We disclosed all external data on the competition forum. Every generated story carries a disclaimer identifying it as AI-generated and inspired by, not representative of, any specific tradition. We also built an evaluation rubric checking for genuine structural authenticity markers, not just the competition's string-similarity metric.

Limitations and uncertainties: This assessment assumes future use remains educational or exploratory rather than commercial or authoritative, an assumption that would need revisiting if the project were deployed more broadly. We cannot fully verify how the model's outputs would be received by the communities whose traditions inspired the supplemental sources, since our validation relied on our own research rather than direct community consultation. This absence of direct community input is itself a limitation worth stating plainly rather than assuming our judgment is sufficient.